

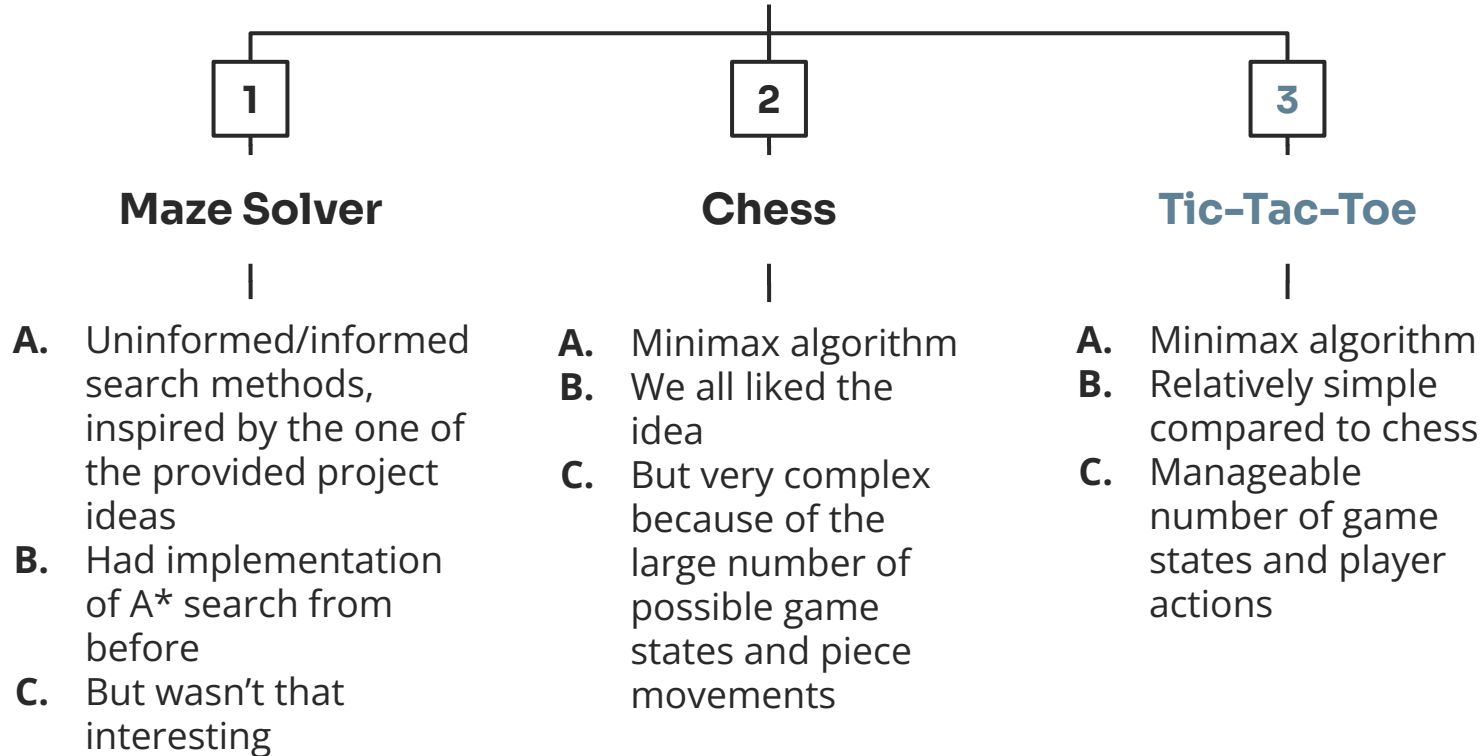


Tic-Tac-Toe AI Solver

Rida Khan, George Mills, Samirah Huda,
Kyle Deng, David George


Initial Goals

AI Game Solver



Objectives

- We aimed to build an AI agent that is able to play Tic-Tac-Toe and win against the opponent
- Our AI agent will use the minimax algorithm to determine the optimal actions to take to win the game
- Our final product will be presented in a GUI
- The user will be able to select X or O, choose the level of difficulty that the AI agent will play by, compete against the agent, and the winner will be displayed on the screen



Methodology:

The Minimax Algorithm

- Used minimax to model and select maximal board states
- Created a score function that would return 100 (X winning final board), 0 (a drawing final board), and -100 (O winning final board)
- Adjusted for depth so that solutions that achieve a winning state in less turns are preferred
- Generated every possible board state from a given board state on every minimax iteration because it's computationally feasible for 3x3 Tic-Tac-Toe
- Informed by *Analysis of Game Tree Search Algorithms Using Minimax Algorithm and Alpha-Beta Pruning* by Shevtekar, Malpe, and Bhaila

Solved Issues

- Traditional minimax implementations is a tree of values, but we needed to return a boardspace and its value
- Generate list of possible immediate moves and corresponding boards, run minimax on each board, compile terminal values in a list, pick min/max, use that index to return future move from list
- Original implementation of minimax resulted in too many optimal states because each of them could eventually result in the winning terminal state
- Make values of win/loss terminal states 100/-100, keep track of depth in minimax, subtract/add final value by depth to prioritize solutions that end the game faster

Methodology:

The Minimax Algorithm for 4x4

- 3x3 algorithm can be used on a 4x4 board without modification
- However, without optimizations, the time it takes to run is too long
- Tested potential optimizations using various combinations of depth limits and heuristics
- Went with a heuristic that randomly chose an adjacent tile to a previously placed letter
- Best solution used this heuristic on the first 3 turns, then traditional minimax to determine the best moves thereafter
- Depth limits appear to be detrimental


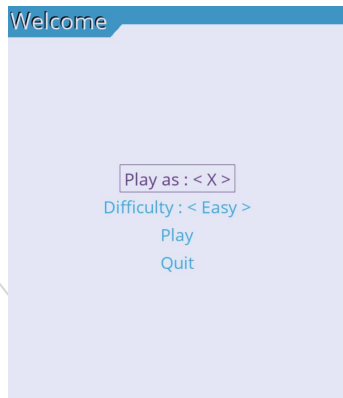
Outcomes

- We were able to build an AI agent that is able to play Tic-Tac-Toe optimally (win or tie) on a 3x3 board using the minimax algorithm
- The user is able to interact with a UI that enables them to choose a letter, the level of difficulty, and compete against the AI agent
- When a player wins, the game ends and the winner is displayed on the screen
- Created another game mode for a 4x4 board in order to better explore the minimax algorithm
 - 4x4 board presented the challenge of balancing depth/breadth limitation for time efficiency and strategy

Areas for Improvement

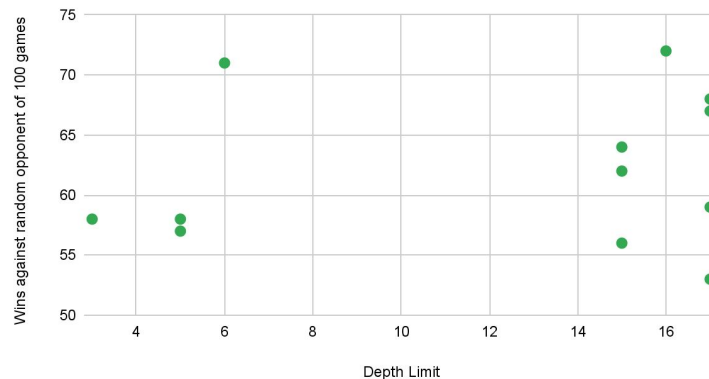
- Accounting for players who do not make optimal moves and make the game more interesting by using expectimax over minimax
- While efficiency for 3x3 is not a concern on our hardware, alpha-beta pruning might make the game more computationally accessible to worse/older hardware
- Implementing more rigorous heuristics, more informed breadth/depth limits, and pruning to improve 4x4

Game Images



Quantitative Metrics

Depth vs Average Win Rate against Random



Breadth vs. Average Win Rate against Random

